

Bắt đầu lập trình hoặc tạo mã bằng trí tuệ nhân tạo (AI).

Nhấp đúp (hoặc nhấn Enter) để chỉnh sửa

```
# === BƯỚC 0: Thiết lập & Tải dữ liệu ===
import os, tarfile
import pandas as pd
from sklearn.preprocessing import LabelEncoder

# 1) Nếu đã có file .tar.gz trong data/, giải nén
TAR_PATH = "data/hwu.tar.gz"
if os.path.exists(TAR_PATH):
    with tarfile.open(TAR_PATH, "r:gz") as tar:
        tar.extractall("data")
    print(" Đã giải nén:", TAR_PATH)
else:
    # 2) Nếu chưa có, cho phép upload trực tiếp 1 file .tar.gz HOẶC 3 file csv (train/val/test)
    try:
        from google.colab import files
        print(" Upload hwu.tar.gz HOẶC train/val/test (.csv)")
        up = files.upload()
        os.makedirs("data", exist_ok=True)
        for name, content in up.items():
            open(os.path.join("data", name), "wb").write(content)
        # Tự giải nén nếu người dùng upload .tar.gz
        for name in up.keys():
            if name.endswith(".tar.gz") or name.endswith(".tgz"):
                with tarfile.open(os.path.join("data", name), "r:gz") as tar:
                    tar.extractall("data")
                print(" Đã giải nén:", name)
    except Exception as e:
        print(" Không chạy trong Colab hoặc upload thất bại:", e)

# 3) Xác định thư mục chứa file sau giải nén
data_dir = "data/hwu" if os.path.isdir("data/hwu") else "data"
print(" data_dir =", data_dir, "| Files:", os.listdir(data_dir))

# 4) Đọc đúng file CSV (theo gói bạn vừa giải nén)
train_path = os.path.join(data_dir, "train.csv")
val_path = os.path.join(data_dir, "val.csv")
test_path = os.path.join(data_dir, "test.csv")
assert os.path.exists(train_path) and os.path.exists(val_path) and os.path.exists(test_path), \
    "Không tìm thấy train.csv/val.csv/test.csv trong " + data_dir

# 5) Đọc dữ liệu (CSV, header chuẩn: text + intent). Nếu header khác, đổi tên cột tương ứng.
df_train = pd.read_csv(train_path)
df_val = pd.read_csv(val_path)
df_test = pd.read_csv(test_path)

# Chuẩn hoá tên cột nếu khác (ví dụ 'sentence', 'query' → 'text'; 'label', 'category' → 'intent')
def normalize_cols(df):
    cols = {c.lower(): c for c in df.columns}
    text_col = None
    for k in ["text", "utterance", "sentence", "query", "content"]:
        if k in cols: text_col = cols[k]; break
    intent_col = None
    for k in ["intent", "label", "category", "class", "target"]:
        if k in cols: intent_col = cols[k]; break
    if text_col is None: text_col = df.columns[0]
    if intent_col is None: intent_col = df.columns[1]
    return df.rename(columns={text_col: "text", intent_col: "intent"})[["text", "intent"]]

df_train = normalize_cols(df_train)
df_val = normalize_cols(df_val)
df_test = normalize_cols(df_test)

print("Train shape:", df_train.shape)
print("Validation shape:", df_val.shape)
print("Test shape:", df_test.shape)
display(df_train.head())

# 6) Mã hoá nhãn bằng LabelEncoder (fit TRÊN CẢ train+val+test để đồng bộ bộ nhãn)
le = LabelEncoder()
```

```
le.fit(pd.concat([df_train["intent"], df_val["intent"], df_test["intent"]], axis=0))

y_train = le.transform(df_train["intent"])
y_val = le.transform(df_val["intent"])
y_test = le.transform(df_test["intent"])
num_classes = len(le.classes_)
print(" num_classes:", num_classes, "| ví dụ nhãn:", list(le.classes_)[:10], "...")
```

Upload hwu.tar.gz HOẶC train/val/test (.csv)

Chọn tệp Chưa có tệp nào được chọn Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving hwu.tar.gz to hwu.tar (6).gz

data_dir = data/hwu | Files: ['train_10.csv', 'categories.json', 'val.csv', 'train_5.csv', 'test.csv', 'train.csv']

Train shape: (8954, 2)

Validation shape: (1076, 2)

Test shape: (1076, 2)

	text	intent
0	what alarms do i have set right now	alarm_query
1	checkout today alarm of meeting	alarm_query
2	report alarm settings	alarm_query
3	see see for me the alarms that you have set to...	alarm_query
4	is there an alarm for ten am	alarm_query

num_classes: 64 | ví dụ nhãn: ['alarm_query', 'alarm_remove', 'alarm_set', 'audio_volume_down', 'audio_volume_mute', 'audio_vol

❏ **Nhiệm vụ 1**

```
# === NHIỆM VỤ 1: TF-IDF + Logistic Regression ===
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.pipeline import make_pipeline
from sklearn.metrics import classification_report, f1_score

tfidf_lr_pipeline = make_pipeline(
    TfidfVectorizer(max_features=5000),
    LogisticRegression(max_iter=1000, n_jobs=-1, random_state=42)
)

tfidf_lr_pipeline.fit(df_train["text"], y_train)
y_pred_lr = tfidf_lr_pipeline.predict(df_test["text"])

print("=== Classification report: TF-IDF + LR ===")
print(classification_report(y_test, y_pred_lr, target_names=le.classes_, digits=4))
f1_lr = f1_score(y_test, y_pred_lr, average="macro")
print("Macro-F1 (test):", f1_lr)
```

email_addcontact	0.7778	0.8750	0.8235	8
email_query	0.8333	0.7895	0.8108	19
email_querycontact	0.9231	0.6316	0.7500	19
email_sendemail	0.8095	0.8947	0.8500	19
general_affirm	1.0000	1.0000	1.0000	19
general_commandstop	1.0000	1.0000	1.0000	19
general_confirm	1.0000	1.0000	1.0000	19
general_dontcare	0.9048	1.0000	0.9500	19
general_explain	1.0000	0.9474	0.9730	19
general_joke	1.0000	1.0000	1.0000	12
general_negate	0.9500	1.0000	0.9744	19

music_likeness	0.6471	0.6111	0.6286	18
music_query	0.7143	0.5263	0.6061	19
music_settings	1.0000	0.5714	0.7273	7
news_query	0.7500	0.6316	0.6857	19
play_audiobook	0.9474	0.9474	0.9474	19
play_game	0.8125	0.6842	0.7429	19
play_music	0.5833	0.7368	0.6512	19
play_podcasts	1.0000	0.8421	0.9143	19
play_radio	0.8889	0.8421	0.8649	19
qa_currency	0.9444	0.8947	0.9189	19
qa_definition	0.8182	0.9474	0.8780	19
qa_factoid	0.4783	0.5789	0.5238	19
qa_maths	0.9231	0.8571	0.8889	14
qa_stock	1.0000	0.9474	0.9730	19
recommendation_events	0.8333	0.7895	0.8108	19
recommendation_locations	0.8095	0.8947	0.8500	19
recommendation_movies	1.0000	1.0000	1.0000	10
social_post	0.9500	1.0000	0.9744	19
social_query	0.8000	0.8889	0.8421	18
takeaway_order	0.8333	0.7895	0.8108	19
takeaway_query	0.8947	0.8947	0.8947	19
transport_query	0.6818	0.7895	0.7317	19
transport_taxi	1.0000	1.0000	1.0000	18
transport_ticket	0.9375	0.7895	0.8571	19
transport_traffic	1.0000	0.9474	0.9730	19
weather_query	0.6190	0.6842	0.6500	19
accuracy			0.8355	1076
macro avg	0.8452	0.8343	0.8353	1076
weighted avg	0.8422	0.8355	0.8351	1076

Macro-F1 (test): 0.8352983005857358

~ Nhiệm vụ 2

```
# === NHIỆM VỤ 2: Word2Vec Avg + Dense ===
import numpy as np
from gensim.models import Word2Vec
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.callbacks import EarlyStopping

# 1) Huấn luyện Word2Vec trên văn bản train
sentences = [str(s).split() for s in df_train["text"]]
w2v_dim = 100
w2v = Word2Vec(sentences=sentences, vector_size=w2v_dim, window=5, min_count=1, workers=4, seed=42)

# 2) Hàm chuyển câu -> vector trung bình
def sentence_to_avg_vector(text, model, dim=100):
    toks = str(text).split()
    vecs = [model.wv[t] for t in toks if t in model.wv]
    if not vecs:
        return np.zeros(dim, dtype="float32")
    return np.mean(vecs, axis=0).astype("float32")

def to_matrix(texts, model, dim=100):
    return np.vstack([sentence_to_avg_vector(t, model, dim) for t in texts])

X_train_avg = to_matrix(df_train["text"], w2v, w2v_dim)
X_val_avg = to_matrix(df_val["text"], w2v, w2v_dim)
X_test_avg = to_matrix(df_test["text"], w2v, w2v_dim)

# 3) Mô hình Dense
model_avg = Sequential([
    Dense(128, activation='relu', input_shape=(w2v_dim,)),
    Dropout(0.5),
    Dense(num_classes, activation='softmax')
])
model_avg.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

es = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True, verbose=0)
_ = model_avg.fit(X_train_avg, y_train, validation_data=(X_val_avg, y_val),
    epochs=50, batch_size=64, callbacks=[es], verbose=0)

y_pred_avg = np.argmax(model_avg.predict(X_test_avg, verbose=0), axis=1)
print("=== Classification report: W2V-Avg + Dense ===")
print(classification_report(y_test, y_pred_avg, target_names=le.classes_, digits=4))
```

```
from sklearn.metrics import f1_score
f1_avg = f1_score(y_test, y_pred_avg, average="macro")
print("Macro-F1 (test):", f1_avg)
```

general_dontcare	0.2037	0.5789	0.3014	19
general_explain	0.0769	0.1053	0.0889	19
general_joke	0.1667	0.0833	0.1111	12
general_negate	0.4615	0.3158	0.3750	19
general_praise	0.3529	0.6316	0.4528	19
general_quirky	0.0000	0.0000	0.0000	19
general_repeat	0.3333	0.3158	0.3243	19
iot_cleaning	0.4375	0.4375	0.4375	16
iot_coffee	0.3810	0.4211	0.4000	19
iot_hue_lightchange	0.4091	0.4737	0.4390	19
iot_hue_lightdim	0.4000	0.1667	0.2353	12
iot_hue_lightoff	0.3208	0.8947	0.4722	19
iot_hue_lighton	0.0000	0.0000	0.0000	3
iot_hue_lightup	0.0000	0.0000	0.0000	14
iot_wemo_off	0.0000	0.0000	0.0000	9
iot_wemo_on	0.3333	0.2857	0.3077	7
lists_createoradd	0.2143	0.4737	0.2951	19
lists_query	0.0000	0.0000	0.0000	19
lists_remove	0.2963	0.4211	0.3478	19
music_likeness	0.0000	0.0000	0.0000	18
music_query	0.1667	0.1053	0.1290	19
music_settings	0.0000	0.0000	0.0000	7
news_query	0.0870	0.1053	0.0952	19
play_audiobook	0.0833	0.0526	0.0645	19
play_game	0.1538	0.1053	0.1250	19
play_music	0.0000	0.0000	0.0000	19
play_podcasts	0.0000	0.0000	0.0000	19
play_radio	0.0769	0.0526	0.0625	19
qa_currency	0.2000	0.0526	0.0833	19
qa_definition	0.0000	0.0000	0.0000	19
qa_factoid	0.0857	0.1579	0.1111	19
qa_maths	0.0000	0.0000	0.0000	14
qa_stock	0.0000	0.0000	0.0000	19
recommendation_events	0.1607	0.4737	0.2400	19
recommendation_locations	0.0000	0.0000	0.0000	19
recommendation_movies	0.0000	0.0000	0.0000	10
social_post	0.0000	0.0000	0.0000	19
social_query	0.0000	0.0000	0.0000	18
takeaway_order	0.1667	0.1053	0.1290	19
takeaway_query	0.5000	0.0526	0.0952	19
transport_query	0.0000	0.0000	0.0000	19
transport_taxi	0.0000	0.0000	0.0000	18
transport_ticket	0.1690	0.6316	0.2667	19
transport_traffic	0.0000	0.0000	0.0000	19
weather_query	0.0000	0.0000	0.0000	19
accuracy			0.2035	1076
macro_avg	0.1667	0.1891	0.1490	1076
weighted_avg	0.1646	0.2035	0.1557	1076

Macro-F1 (test): 0.1490183362996796

/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined for classes in labels [0]: no non-zero entries found in y_true
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined for classes in labels [0]: no non-zero entries found in y_true
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined for classes in labels [0]: no non-zero entries found in y_true
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

▼ Nhiệm vụ 3

```
# === NHIỆM VỤ 3: Embedding (pre-trained) + LSTM ===
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.layers import Embedding, LSTM
from tensorflow.keras.models import Sequential
```

```
# 1) Tokenizer + Padding
max_words = 20000
oov_tok = "<UNK>"
tokenizer = Tokenizer(num_words=max_words, oov_token=oov_tok)
tokenizer.fit_on_texts(df_train["text"])
```

```
def to_pad(texts, tok, max_len=50):
    seqs = tok.texts_to_sequences(texts)
```

```

return pad_sequences(seqs, maxlen=max_len, padding='post', truncating='post')

max_len = 50
X_train_pad = to_pad(df_train["text"], tokenizer, max_len)
X_val_pad = to_pad(df_val["text"], tokenizer, max_len)
X_test_pad = to_pad(df_test["text"], tokenizer, max_len)

vocab_size = min(max_words, len(tokenizer.word_index) + 1)
embedding_dim = w2v_dim

# 2) Ma trận embedding từ W2V (đã train ở Nhiệm vụ 2)
embedding_matrix = np.zeros((vocab_size, embedding_dim), dtype="float32")
for word, idx in tokenizer.word_index.items():
    if idx < vocab_size and word in w2v.wv:
        embedding_matrix[idx] = w2v.wv[word]

# 3) LSTM với embedding pre-trained (đóng băng)
lstm_pre = Sequential([
    Embedding(input_dim=vocab_size, output_dim=embedding_dim,
              weights=[embedding_matrix], input_length=max_len, trainable=False),
    LSTM(128, dropout=0.2, recurrent_dropout=0.2),
    Dense(num_classes, activation='softmax')
])
lstm_pre.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

es = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True, verbose=0)
_ = lstm_pre.fit(X_train_pad, y_train, validation_data=(X_val_pad, y_val),
                epochs=30, batch_size=64, callbacks=[es], verbose=0)

y_pred_pre = np.argmax(lstm_pre.predict(X_test_pad, verbose=0), axis=1)
print("=== Classification report: Emb(pretrained) + LSTM ===")
print(classification_report(y_test, y_pred_pre, target_names=le.classes_, digits=4))
f1_pre = f1_score(y_test, y_pred_pre, average="macro")
print("Macro-F1 (test):", f1_pre)

```

general_dontcare	0.0306	0.3684	0.0565	19
general_explain	0.0000	0.0000	0.0000	19
general_joke	0.0000	0.0000	0.0000	12
general_negate	0.0820	0.2632	0.1250	19
general_praise	0.1053	0.6316	0.1805	19
general_quirky	0.0000	0.0000	0.0000	19
general_repeat	0.1600	0.4211	0.2319	19
iot_cleaning	0.2222	0.2500	0.2353	16
iot_coffee	0.0000	0.0000	0.0000	19
iot_hue_lightchange	0.2558	0.5789	0.3548	19
iot_hue_lightdim	0.0000	0.0000	0.0000	12
iot_hue_lightoff	0.4444	0.8421	0.5818	19
iot_hue_lighton	0.0000	0.0000	0.0000	3
iot_hue_lightup	0.0000	0.0000	0.0000	14
iot_wemo_off	0.0000	0.0000	0.0000	9
iot_wemo_on	0.0000	0.0000	0.0000	7
lists_createoradd	0.0129	0.1053	0.0230	19
lists_query	0.0000	0.0000	0.0000	19
lists_remove	0.0000	0.0000	0.0000	19
music_likeness	0.0000	0.0000	0.0000	18
music_query	0.0357	0.0526	0.0426	19

accuracy			0.0818	1076
macro avg	0.0266	0.0730	0.0362	1076
weighted avg	0.0294	0.0818	0.0402	1076

Macro-F1 (test): 0.03616698778099441

```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined for classes in labels [1]: no samples for any class in the true labels
_warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined for classes in labels [1]: no samples for any class in the true labels
_warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined for classes in labels [1]: no samples for any class in the true labels
_warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
```

▼ Nhiệm vụ 4

```
# === NHIỆM VỤ 4: Embedding (scratch) + LSTM ===
lstm_scr = Sequential([
    Embedding(input_dim=vocab_size, output_dim=100, input_length=max_len),
    LSTM(128, dropout=0.2, recurrent_dropout=0.2),
    Dense(num_classes, activation='softmax')
])
lstm_scr.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

es = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True, verbose=0)
_ = lstm_scr.fit(X_train_pad, y_train, validation_data=(X_val_pad, y_val),
                epochs=30, batch_size=64, callbacks=[es], verbose=0)

y_pred_scr = np.argmax(lstm_scr.predict(X_test_pad, verbose=0), axis=1)
print("=== Classification report: Emb(scratch) + LSTM ===")
print(classification_report(y_test, y_pred_scr, target_names=le.classes_, digits=4))
f1_scr = f1_score(y_test, y_pred_scr, average="macro")
print("Macro-F1 (test):", f1_scr)
```

general_dontcare	0.4412	0.7895	0.5660	19
general_explain	0.5185	0.7368	0.6087	19
general_joke	0.0000	0.0000	0.0000	12
general_negate	0.0000	0.0000	0.0000	19
general_praise	0.3333	0.4211	0.3721	19
general_quirky	0.0000	0.0000	0.0000	19
general_repeat	0.1250	0.0526	0.0741	19
iot_cleaning	0.0000	0.0000	0.0000	16
iot_coffee	0.2759	0.4211	0.3333	19
iot_hue_lightchange	0.1923	0.5263	0.2817	19
iot_hue_lightdim	0.0000	0.0000	0.0000	12
iot_hue_lightoff	0.0000	0.0000	0.0000	19
iot_hue_lighton	0.0000	0.0000	0.0000	3
iot_hue_lightup	0.0000	0.0000	0.0000	14
iot_wemo_off	0.0000	0.0000	0.0000	9
iot_wemo_on	0.0000	0.0000	0.0000	7
lists_createoradd	0.1923	0.5263	0.2817	19
lists_query	0.0000	0.0000	0.0000	19
lists_remove	0.2963	0.4211	0.3478	19
music_likeness	0.0000	0.0000	0.0000	18
music_query	0.1250	0.1579	0.1395	19

accuracy			0.2314	1076
macro avg	0.1293	0.2049	0.1495	1076
weighted avg	0.1458	0.2314	0.1688	1076

Macro-F1 (test): 0.14952501738646845

```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined for classes in true labels but not predicted: no samples predicted on class 0. Precision will be set to nan.
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined for classes in true labels but not predicted: no samples predicted on class 0. Precision will be set to nan.
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined for classes in true labels but not predicted: no samples predicted on class 0. Precision will be set to nan.
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

▼ Nhiệm vụ 5

```
# === NHIỆM VỤ 5: So sánh định lượng + Phân tích định tính ===
import numpy as np
import pandas as pd
from sklearn.metrics import f1_score, classification_report

# ----- 1 So sánh định lượng (bảng tổng hợp F1 và Loss) -----
loss_avg, _ = model_avg.evaluate(X_test_avg, y_test, verbose=0)
loss_pre, _ = lstm_pre.evaluate(X_test_pad, y_test, verbose=0)
loss_scr, _ = lstm_scr.evaluate(X_test_pad, y_test, verbose=0)

summary = pd.DataFrame({
    "Pipeline": [
        "TF-IDF + Logistic Regression",
        "Word2Vec (Avg) + Dense",
        "Embedding (Pre-trained) + LSTM",
        "Embedding (Scratch) + LSTM"
    ],
    "F1-macro (test)": [f1_lr, f1_avg, f1_pre, f1_scr],
    "Test Loss": [None, loss_avg, loss_pre, loss_scr]
}).sort_values("F1-macro (test)", ascending=False).reset_index(drop=True)

display(summary)

# ----- 2 Phân tích định tính: các câu "khó" -----
hard_texts = [
    "can you remind me to not call my mom",          # phủ định
    "is it going to be sunny or rainy tomorrow",     # lựa chọn
    "find a flight from new york to london but not through paris" # phủ định + điều kiện
]

# Dự đoán từ 4 mô hình
def predict_all(texts):
    pred_lr = tfidf_lr_pipeline.predict(texts)
    Xavg = np.vstack([sentence_to_avg_vector(t, w2v, w2v_dim) for t in texts])
    pred_avg = np.argmax(model_avg.predict(Xavg, verbose=0), axis=1)
    seqs = tokenizer.texts_to_sequences(texts)
    Xpad = pad_sequences(seqs, maxlen=max_len, padding='post', truncating='post')
    pred_pre = np.argmax(lstm_pre.predict(Xpad, verbose=0), axis=1)
    pred_scr = np.argmax(lstm_scr.predict(Xpad, verbose=0), axis=1)
    return pred_lr, pred_avg, pred_pre, pred_scr

pred_lr, pred_avg, pred_pre, pred_scr = predict_all(hard_texts)

# Lấy nhãn thật (nếu có trong test set, nếu không sẽ ghi "N/A")
true_labels = []
for t in hard_texts:
    match = df_test[df_test["text"].str.lower() == t.lower()]
    if len(match) > 0:
        true_labels.append(le.transform(match["intent"])[0])
    else:
        true_labels.append(None)

# Bảng kết quả định tính
rows = []
for i, text in enumerate(hard_texts):
    y_true = true_labels[i]
    row = {
        "text": text,
        "True Intent": le.classes_[y_true] if y_true is not None else "N/A",
        "TF-IDF + LR": le.classes_[pred_lr[i]],
        "W2V-Avg + Dense": le.classes_[pred_avg[i]]
    }
    rows.append(row)
```

```

    "LSTM (pre)": le.classes_[pred_pre[i]],
    "LSTM (scratch)": le.classes_[pred_scr[i]],
}
# Thêm dấu ✓ nếu có nhãn thật và dự đoán đúng
if y_true is not None:
    row["✓ LR"] = "✓" if pred_lr[i] == y_true else "X"
    row["✓ W2V"] = "✓" if pred_avg[i] == y_true else "X"
    row["✓ LSTM-pre"] = "✓" if pred_pre[i] == y_true else "X"
    row["✓ LSTM-scr"] = "✓" if pred_scr[i] == y_true else "X"
rows.append(row)

qual_df = pd.DataFrame(rows)
display(qual_df)

# ----- 3 Tổng kết nhận xét -----
print("\n Nhận xét:")
print("- Các câu có phủ định ('not') hoặc mệnh đề phụ ('or', 'but not through ...') thường khiến mô hình TF-IDF và W2V trung bình
    "vì chúng không nắm được thứ tự và phạm vi của phủ định.")
print("- LSTM (đặc biệt bản pre-trained) có xu hướng hiểu ngữ cảnh tốt hơn, vì trạng thái ẩn theo chuỗi giúp mô hình nhận diện đ
print("- Nếu LSTM (pre-trained) chính xác hơn bản học từ đầu, điều đó chứng tỏ embedding Word2Vec cung cấp ngữ nghĩa ban đầu hữu
```

	Pipeline	F1-macro (test)	Test Loss				
0	TF-IDF + Logistic Regression	0.835298	NaN	 			
1	Embedding (Scratch) + LSTM	0.149525	2.941058				
2	Word2Vec (Avg) + Dense	0.149018	3.032118				
3	Embedding (Pre-trained) + LSTM	0.036167	3.603505				
	text	True Intent	TF-IDF + LR	W2V-Avg + Dense	LSTM (pre)	LSTM (scratch)	
0	can you remind me to not call my mom	N/A	calendar_set	general_joke	general_dontcare	email_sendemail	
1	is it going to be sunny or rainy tomorrow	N/A	weather_query	calendar_query	lists_createoradd	social_post	
2	find a flight from new york to london but not ...	N/A	general_negate	transport_query	alarm_set	cooking_recipe	

Nhận xét:

- Các câu có phủ định ('not') hoặc mệnh đề phụ ('or', 'but not through ...') thường khiến mô hình TF-IDF và W2V trung bình dự đo
- LSTM (đặc biệt bản pre-trained) có xu hướng hiểu ngữ cảnh tốt hơn, vì trạng thái ẩn theo chuỗi giúp mô hình nhận diện được mỗ
- Nếu LSTM (pre-trained) chính xác hơn bản học từ đầu, điều đó chứng tỏ embedding Word2Vec cung cấp ngữ nghĩa ban đầu hữu ích, g